

# Changing Linux Page Replacement Algorithm

---

Zihan Wang

# Overview

---

- Study of the Original Algorithm
- Changing to an LFU-Based Algorithm
- Test Results

# Study of the Original Algorithm

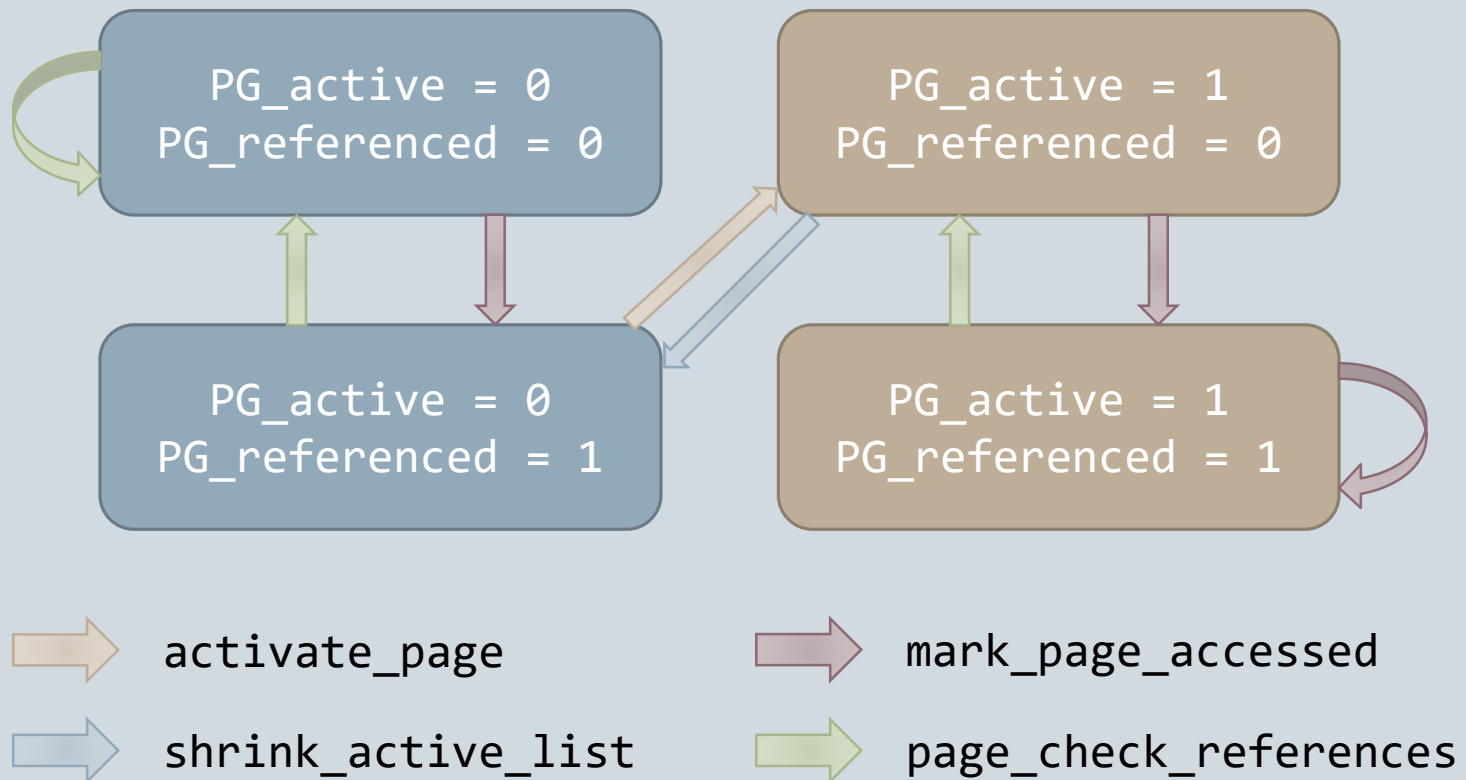
---

# Basic Ideas

---

- Reclaim 'unused' pages only.
- Associate a counter storing the age of the page with each page in RAM.

# State Diagram



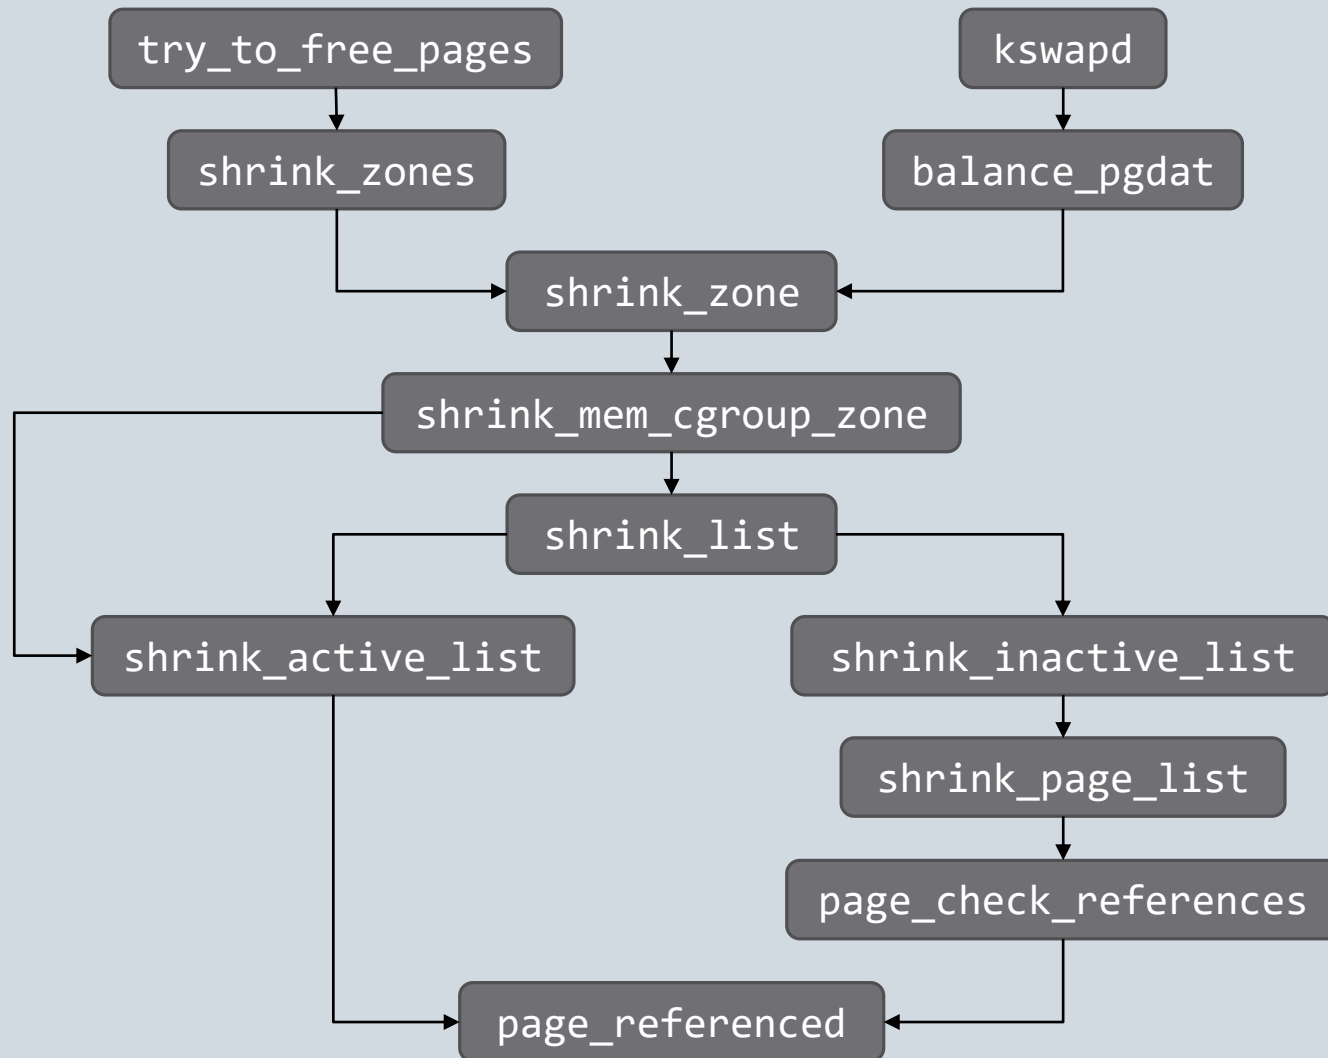
# Condition of Reclaiming

---

- Low on memory: The kernel has difficulty allocating new pages.
- Periodic: Kernel threads are invoked periodically to check whether free pages are above a certain threshold.

## LOW ON MEMORY RECLAIMING

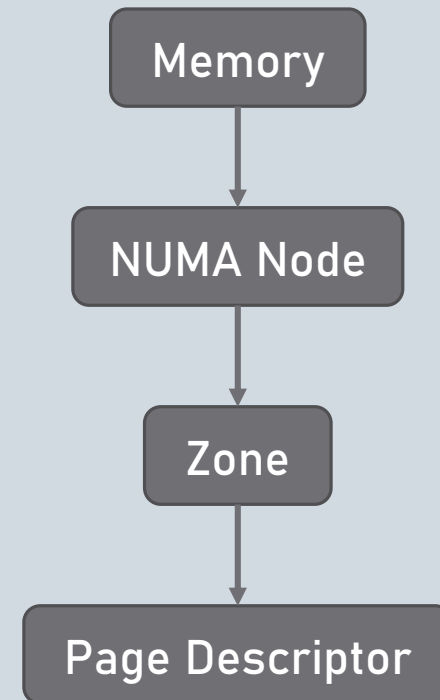
## PERIODIC RECLAIMING



# Memory-Related Structs

---

- NUMA: The time required to access different memory locations is different.
- Memory Zone: Accounting for hardware constraints of the type of data stored in pages.





# Changing to an LFU-Based Algorithm

---

# Scope of Work

---

- The algorithm involved just select the candidates to be replaced.
- We just need to change how the status of the pages are transitioned.

# Comparison of Two Algorithms

Task	Original	LFU
Refresh page	If PG_referenced is 0, set it. Otherwise clear it and add to active list.	Left shift PG_referenced and add by an offset.
Age page	If PG_referenced is 0, clear it. Otherwise set it and add page to inactive list.	Right shift PG_referenced.

Condition	Original	LFU
Add to active list	PG_referenced is 1 and the page is referenced recently.	PG_referenced is above some threshold.
Add to inactive list	PG_referenced is 0 and the page is not referenced recently.	PG_referenced is below some threshold.

# Parameters

---

- The offset to add to PG\_referenced when page is refreshed: 1
- The threshold of moving pages between lists: to be tested
- The maximum limit of PG\_referenced: necessary, also to be tested

# Test Results

---

# Test the Algorithm

---


- Write a program to occupy as much memory as possible.
- Access `/proc/meminfo` periodically to get the sizes of active and inactive lists.

# Notice

---

- kswapd invokes lowmemkiller, which kills processes to reclaim memory space.
- The test program can no longer run while the pages are being reclaimed.

# Variation Trend







# Variation Trend

Time	Trend	Explanation
1-6s	Inactive list continues to grow. Active list remains.	Test program access every page only once.
7-9s	Inactive list grows more slowly. Active list begins to grow.	Page replacement procedures are started.
9-11s	Inactive list grows, while active list shrinks.	Some pages are aged in the replacement algorithm.
11-12s	Inactive list drops dramatically.	Some processes, including the test program, are killed.


# The Impact of Threshold




# The Impact of Threshold



# The Impact of Maximum



# The Impact of Maximum



# Observation

---

- The smaller the threshold, the larger the active list.
- The higher the maximum, the larger the active list.
- The size of active list varies more greatly in the new algorithm, compared with the original.
- The parameters seem to have little impact on the size of inactive list.

# Choice of Parameters

Criterion	Threshold	Maximum
Similar to the original	8	64
Smaller active list size	16 or higher	Threshold $\times$ 2

# Reference

---

- Understanding the Linux Kernel. Daniel Bovet and Marco Cesati. O'Reilly Media.



# Thanks

---